



ForgeMind™

The GIS Intent Compiler

Natural language in. Correct geospatial pipelines out. Or an honest refusal. ForgeMind™ is the natural-language control plane for ForgeGIS™: it turns a plain-language request into a validated, executed pipeline on the GPU-accelerated engine, and when a request cannot be honored correctly it returns an explicit refusal rather than a confident guess. It runs on the AI provider you choose — or on a model you host yourself.

Zero

confident-wrong answers

across the June 2026 conversation validation campaign

Refuse

rather than guess

on the compiled path a request becomes a validated pipeline or a named refusal

Your_{model}

four vendors or your own

Claude, OpenAI, Gemini or Grok — or an OpenAI-compatible server you run, such as Ollama

What it is

ForgeMind sits between a person, or another agent, and the ForgeGIS engine. It decides what to run, binds the request to real data through ForgeData™, dispatches the work and narrates the result — and the answer is drawn on the map you asked from, in ForgeGIS Studio™ or Seaglass Globe™.

Why it matters

A conventional agent lets a model write tool calls freely, so every wrong parameter lands looking confident. On its compiled path ForgeMind narrows the model to understanding intent and lets deterministic code own the pipeline, so what is not sound is declined instead of run. That is what regulated and mission work needs.

How it's built differently

For its templated capabilities the same parsed intent produces the same pipeline every time. With a hosted vendor a fast, inexpensive model handles intent, and since 1.2.5 most answers are narrated from templates with no model call. On one site-suitability request (Claude, May 2026) the compiled path cost about a twenty-fifth of a legacy agent-loop run that never finished.

WHO FORGEMIND IS FOR

Operators and analysts

Ask for a viewshed, a slope mask or a site-suitability search in plain language and get a correct result without learning a desktop GIS. ForgeMind holds context across a conversation — establish a subject once, then refine it: "re-run that with a wider radius", "mask it to slopes under seven degrees", "rebuild it — the objective moved".

Programs that cannot accept a plausible wrong answer

When a request is under-specified, out of scope or cannot be bound to data, ForgeMind asks, says no data covers the area, or refuses with a reason. With a model server on the site's own network, no request goes to a cloud model vendor at all.

Teams building on AI agents

ForgeMind is native to the Model Context Protocol in both directions: it drives the suite's engines as tools, and it can be one — a single tool a higher-level agent hands a whole geospatial workflow to, plus a structured surface that runs a capability with no model call at all.

NEW IN 1.2.5

Your provider, your hardware. Gemini and Grok join Claude and OpenAI, and any OpenAI-compatible server you run plugs into the same client — Ollama is the tested path. The answer lands on the map you asked from, with the scene you were looking at sent along with the question, and with the compute broker installed, a folder batch such as "hillshade every tile in this folder" goes to it as one job.

What the figures rest on. Zero confident-wrong is the June 2026 conversation campaign, run before the 1.2.5 provider options; the model-spend comparison is one request on Claude against a run that hit its iteration cap. The Technical Brief sets out what each figure shows and what it does not, and where refusal by construction stops.

Get the full Technical Brief

The Technical Brief covers the architecture, the providers, validation and its scope, deployment and security. ForgeMind is part of the Forge suite, installed by the Seaglass Control Center™. The Forge suite is commercially licensed; evaluation builds are available to qualified teams.

rich@seaglassfoundry.com · seaglassfoundry.com