SEAGLASS FOUNDRY

Forged For Speed. Refined by Craft.

ForgeMind™

The GIS Intent Compiler

Natural language in. Correct geospatial pipelines out. Or an honest refusal. ForgeMind™ is the natural-language control plane for ForgeGIS™: it turns a plain-language request into a validated, executed pipeline on the GPU-accelerated engine, and when a request cannot be honored correctly it returns an explicit refusal rather than a confident guess. It runs on the AI provider you choose — or on a model you host yourself.

Zero

confident-wrong answers

across the June 2026 conversation validation campaign

Refuse

rather than guess

on the compiled path a request becomes a validated pipeline or a named refusal

Your model

four vendors or your own

Claude, OpenAI, Gemini or Grok — or an OpenAI-compatible server you run, such as Ollama

What it is

ForgeMind sits between a person, or another agent, and the ForgeGIS engine. It decides what to run, binds the request to real data through ForgeData™, dispatches the work and narrates the result — and the answer is drawn on the map you asked from, in ForgeGIS Studio™ or Seaglass Globe™.

Why it matters

A conventional agent lets a model write tool calls freely, so every wrong parameter lands looking confident. On its compiled path ForgeMind narrows the model to understanding intent and lets deterministic code own the pipeline, so what is not sound is declined instead of run. That is what regulated and mission work needs.

How it's built differently

For its templated capabilities the same parsed intent produces the same pipeline every time. With a hosted vendor a fast, inexpensive model handles intent, and since 1.2.5 most answers are narrated from templates with no model call. On one site-suitability request (Claude, May 2026) the compiled path cost about a twenty-fifth of a legacy agent-loop run that never finished.

New in 1.2.6 — measured right, and checked. Every distance, area and bearing is measured on the WGS84 ellipsoid; a line between two places is drawn as the geodesic — great circle and rhumb line both, for a navigator's question — and a radius question answers only from inside the radius. Ask "what data do you have for Denver?" and the answer comes from the catalog. Any figure a model writes is checked against what was measured.

CONTINUED ON REVERSE →

WHO FORGEMIND IS FOR

Operators and analysts

Ask for a viewshed, a slope mask or a site-suitability search in plain language and get a correct result without learning a desktop GIS. ForgeMind holds context across a conversation — establish a subject once, then refine it: "re-run that with a wider radius", "mask it to slopes under seven degrees", "rebuild it — the objective moved".

Programs that cannot accept a plausible wrong answer

When a request is under-specified, out of scope or cannot be bound to data, ForgeMind asks, says no data covers the area, or refuses with a reason. Since 1.2.6, a question for a figure that ForgeMind cannot answer gets a reply saying what is missing, never a fallback agent's guess. With a model server on the site's own network, no request goes to a cloud model vendor at all.

Teams building on AI agents

ForgeMind is native to the Model Context Protocol in both directions: it drives the suite's engines as tools, and it can be one — a single tool a higher-level agent hands a whole geospatial workflow to, plus a structured surface that runs a capability with no model call at all.

Anyone who needs to know how sure an answer is

Statistics over a shape on the map are the shape's, not its bounding box's. An answer computed over an area the model placed from memory, or at a place ForgeMind cannot check, says so plainly, and a highest point found outside the radius you asked about is refused rather than reported.

> **Your provider, your hardware.** Claude, OpenAI, Gemini and Grok are built in, and any OpenAI-compatible server you run plugs into the same client — Ollama is the tested path. The answer lands on the map you asked from, with the scene you were looking at sent along with the question, and with the compute broker installed, a folder batch such as "hillshade every tile in this folder" goes to it as one job.

> **What the figures rest on.** Zero confident-wrong is the June 2026 conversation campaign, run before the 1.2.5 provider options; the model-spend comparison is one request on Claude against a run that hit its iteration cap. The Technical Brief sets out what each figure shows and what it does not, and where refusal by construction stops.

Get the full Technical Brief

The Technical Brief covers the architecture, the providers, validation and its scope, deployment and security. ForgeMind is part of the Forge suite, installed by the Seaglass Control Center™. The Forge suite is commercially licensed; evaluation builds are available to qualified teams.

rich@seaglassfoundry.com · seaglassfoundry.com

ForgeMind™, ForgeGIS™, ForgeData™, ForgeGIS Studio™, Seaglass Globe™, Seaglass Control Center™, and Seaglass Foundry™ are trademarks of Seaglass Foundry LLC. Claude is a trademark of Anthropic, PBC; OpenAI is a trademark of OpenAI; Gemini is a trademark of Google LLC; Grok is a trademark of xAI; Ollama and the Model Context Protocol are marks of their respective owners. Reference to these marks does not imply endorsement. © 2026 Seaglass Foundry LLC. All rights reserved.